

Using KVM as a Real-Time Hypervisor


Jan Kiszka, Siemens AG, CT T DE IT 1
Corporate Competence Center Embedded Linux
jan.kiszka@siemens.com


Agenda

- **Motivation & scenarios**
- **RT benchmark updates**
- **Improving QEMU RT performance**
 - Analysis of critical paths
 - Steps to overcome latency spots
- **Summary**

Recall Last Year: Why Using KVM in Embedded?

“We just need a tiny hypervisor to fully exploit this multicore CPU”

- “A few thousand” lines of hypervisor code
- Minimal hardware emulation
- “A bit” paravirtualization 
- Devices are passed through




Recall Last Year: Why Using KVM in Embedded?

“We just need a tiny hypervisor to fully exploit this multicore CPU”

- “A few thousand” lines of hypervisor code
- Minimal hardware emulation
- “A bit” paravirtualization
- Devices are passed through

“But it would be nice to...”

- share some devices
- run upstream Linux and latest Windows




Recall Last Year: Why Using KVM in Embedded?

“We just need a tiny hypervisor to fully exploit this multicore CPU”

- “A few thousand” lines of hypervisor code
- Minimal hardware emulation
- “A bit” paravirtualization
- Devices are passed through

“But it would be nice to...”

- share some devices
- run upstream Linux and latest Windows
- over-commit resources
- manage power




Recall Last Year: Why Using KVM in Embedded?

“We just need a tiny hypervisor to fully exploit this multicore CPU”

- “A few thousand” lines of hypervisor code
- Minimal hardware emulation
- “A bit” paravirtualization
- Devices are passed through

“But it would be nice to...”

- share some devices
- run upstream Linux and latest Windows
- over-commit resources
- manage power
- backup / migrate guests
- use advanced HA features
- ...



...and in Real-Time Scenarios? Pros & Cons

From partitioning hypervisors...

- + High degree of temporal isolation
- + Static allocations simplify RT guarantees
- Poor flexibility
- Non-commodity setup

...and in Real-Time Scenarios? Pros & Cons

From partitioning hypervisors...

- + High degree of temporal isolation
- + Static allocations simplify RT guarantees
- Poor flexibility
- Non-commodity setup

... to full virtualization

- Usually not designed for RT
- Higher complexity makes establishing RT harder
- + Benefit from large user base
 - Guest support
 - Test coverage
- + Benefit from advanced virtualization features
- + RT and SMP scalability share many requirements

Typical Real-Time Guest Setups

Guest types

- Classic RTOS
- AMP (RTOS + x)
- GPOS with RT requirements

Guest interacts with real world – in real-time

- Real-time network (normal/RT Ethernet, fieldbuses, etc.)
- Digital & analogue I/O interfaces
- Data acquisition adapters

Interface access

- Pass-through, i.e. 1:1 mapping of periphery to guest
- Emulation
 - Decoupling of guest driver and host hardware
 - Physical interface sharing – or avoiding (test environments)

What is possible today?

Timed Task Benchmarks: Setup (1)

Host system


- Intel Core i7, 4 cores, 2 threads each
- OpenSUSE 11.4
- **PREEMPT-RT** kernel 2.6.33.9-rt31
- **cyclicttest measures timed task wakeup latency**
`cyclicttest -n -p 99 -h 500 -q`
- Host-side load
 - Cache benchmark loop
`calibrator 3392 8M outputfile`
 - I/O benchmark loop
`echo 1 > /proc/sys/vm/drop_caches ; bonnie -y -s 2000`
- Load loops and cyclicttest (for host benchmark) or guest VCPU thread (for guest benchmark) bound to host CPU 0

Timed Task Benchmarks: Setup (2)

Guest system


- OpenSUSE 11.4
- PREEMPT-RT kernel 2.6.33.9-rt31
- qemu-kvm patched to allow **prioritization**
- VM configured to avoid latency-sensitive guest exits:
`-m 1G -drive file=guest.img,if=virtio`
`-rt maxprio=80,paioprio=1 -nographic -vga none`
`-netdev user,hostfwd=::2222-:22,id=net`
`-net nic,netdev=net`
- cyclicttest measures timed guest task wakeup latency
`cyclicttest -n -p 99 -h 500 -q`
- Host-side load kept unchanged

Timed Task Benchmarks: Results after ~3h



cyclicttest on host
Maximum: 29 μ s

cyclicttest on guest
Maximum: 112 μ s



Note: Test length too short for reliable maxima

External Event Benchmark: AMP RT Guest with Passed-Through NIC

Host configuration

- Base setup as before
- Intel i82541PI NIC as I/O device (no MSI)
- VM with 2 VCPUs

Guest properties

- GPOS and RTOS on different VCPUs
- RTOS only interacts with
 - APIC & IO-APIC
 - Assigned devices (here: PCI NIC)**=> no exits to QEMU user space**
- GPOS requires full-blown virtualization, specifically VGA

External Event Benchmark: Measuring Network Latency


External measurement system

- Linux/Xenomai with RTnet
- rtping @100 HZ

Load scenario

- `hackbench 150 process 1000`
- Disk I/O load on host
- `ping -f` from host to GPOS guest (via tap+virtio)
- ftrace enabled for events

Worst case round-trip latency:
(after 16 h)



330 μs

External Event Benchmark: Measuring Network Latency

External measurement system


- Linux/Xenomai with RTnet
- rtping @100 HZ

Load scenario

- `hackbench 150 process 1000`
- Disk I/O load on host
- `ping -f` from host to GPOS guest (via tap+virtio)
- ftrace enabled for events

Worst case round-trip latency:
(after 16 h)

Same scenario with emulated NIC:
(prioritized host NIC IRQ & RX Soft IRQ)



330 μs

100 ms – and more

QEMU Still Ruining Latencies

Everything under `qemu_global_mutex`

- Remaining synchronous disk I/O

*Note: observed `io_submit()` syscall latencies >1 s,
paio architecture is immune*

- Network I/O
- Terminal I/O
- X interaction (GUI updates)
- Dirty RAM log synchronization
(>10 ms on `synchronize_srcu_expedited`)
- ...and probably more

`qemu_global_mutex` is a no-go for RT code paths!

Overcoming the Global Lock – Road Works

CPUState

- Read/write access
- `cpu_single_env`

PIO/MMIO request-to-device dispatching

Coalesced MMIO flushing

Back-end access

- TX on network layer
- Write to character device
- Timer setup, etc.

Back-end events (iothread jobs)

- Network RX, read from chardev, timer signals, ...

IRQ delivery

- Raising/lowering from device model to IRQ chip
- Injection into VCPU (if user space IRQ chips)

Step 1: Localize CPUState

VCPU owns its CPUState

- No remote write unless VCPU is stopped
- Establish formal rule
(pre-exists for KVM core)
- Just few code changes required

cpu_current_env becomes per-CPU variable

- pthread_set/get_specific on UNIX
- Win32 requires wrapping
- Works with single TCG CPU thread as well

Step 2: I/O Dispatching

Which device handles accessed memory region?

Critical path

- Walk memory map
- Obtain handler & device reference
- Invoke handler
- Done

Preferred approach: lock-less

- Modifications are rare
- Acquiring read-side lock is costly, may even deadlock

Solution: stop machine while modifying memory map
(pattern also used in kvm-tool)

Step 3: Coalesced MMIO Handling

Coalesced MMIO ring as contention point

- One ring per-VM
- Readers must synchronize
- Currently protected by `qemu_global_mutex`

Short-term solution

- Skip flush if target device does not use coalesced MMIO
- Affects VGA and E1000 so far

Long-term solution

- One ring per-device – or MMIO region
- Socket-based `ioeventfd` may be the answer

Step 4: IRQ Forwarding

Typical IRQ path

- Device changes level / generates edge
 - IRQ routers (PCI host, bridges, IRQ remapper, etc.) forward to interrupt controller
 - Interrupt controller forwards to CPU
- => Routing involves multiple device models, i.e. potentially multiple critical sections**

Cannot take the long road if source & sink are in-kernel

- Hacks exist to explore and monitor routes – on x86
- => Generic mechanism required**

Fast path from device to target CPU

- No interaction with routing devices
- State changes (reroutes, blockings) reported to subscribers
- Routing device states can be updated on demand

The Harder Nuts –

Step 5: Concurrent Device Models

Mandatory

- Separate contexts to handle host-originated events
- Enables event prioritization and parallelizing
- iothread(s) can remain “best effort” zone(s)

Variant A

- Per-device lock for atomic sections
- Separate iothreads

Variant B

- Device server thread executes atomic sections

Variant A: A Lock for Every Device

Per-device lock

- Protects atomic sections (PIO/MMIO requests, event processing)
- Can be taken over VCPU or I/O thread contexts

Separate I/O threads

- Process host-triggered work
 - Device-related file descriptor callbacks
 - Bottom-halves
- Granularity: device or group of devices

Downside

- MMIO addresses device, device issues DMA to another device
=> lock nestings, lock recursions, deadlocks
- Which lock to acquire in which order?
- Can we drop the device lock while calling core services?

Variant B: Device Server Thread

Server thread runs device jobs

- Host-triggered work
- Complex guest-triggered work

Guest I/O requests forwarded to server

- Write requests can be synchronous and asynchronous
- Reads must be synchronous

Trivial I/O requests do not require server context

- get/set register without side effects

Thread ensures atomicity of device model

=> **no locks required** (famous last words...)

Downsides

- May require careful ordering of state changes
- May require use of atomics & barriers

Work in Progress

QEMU activities

- Implement sketched road map
- Currently focusing on variant B
- Primary target
 - E1000 device model
 - KVM with in-kernel IRQ chips

Kernel activities

- Hunt & analyze potential latency spots (hundred μ s range)
- Address IRQ thread management issue

Implementation Footnote: Fun with glibc and POSIX

glibc's condition variables
+ priority inheritance mutexes
= deadlock

Background

- Internal condvar locks aren't PI-aware
- Using PI locks unconditionally considered too heavy
- Lacking POSIX interface to declare PI for condvars
- Patches exist for pthread_condattr_setprotocol_np
- Ignored by glibc folks :-)

Workarounds

- Use priority ceiling
 - Costly (one syscall per mutex lock/unlock)
 - All participating threads must be SCHED_FIFO/RR
- Don't use condvars

Summary

Many benefits of using KVM as RT hypervisor

- Full virtualization feature set
- Matured support for broad range of guests

Restricted RT support so far

- Works well without QEMU in the loop
- User space VM exits trigger huge latencies

Ongoing work to reduce restrictions

- Parallelize and prioritize QEMU device models
- Next goal: emulated RT networking
- Event loop latencies $\ll 1$ ms in reach

Progress on real-time will improve SMP scalability as well!

Any Questions?

Thank you!